Instruction Level Parallelism

Tricking the system to maximize throughput

Source material

- D. A. Patterson and J. L. Hennessy, *Computer Organization and Design: The Hardware/Software Interface: RISC-V Edition*. Cambridge, MA: Morgan Kaufmann, 2018.
    - Chapter 4.7: Data Hazards: Forwarding versus Stalling
    - Chapter 4.8: Control Hazards
    - Chapter 4.10: Parallelism via Instructions
- J. L. Hennessy, D. A. Patterson, and K. Asanović, *Computer Architecture: A Quantitative Approach*, 5th ed. Waltham, MA: Morgan Kaufmann/Elsevier, 2012.
    - Appendix C: Pipelining: Basic and Intermediate Concepts
    - Chapter 3: Instruction-Level Parallelism and Its Exploitation

Goal: Improving throughput by increasing ILP

1. Revisit how we circumvent hazards to avoid pipeline stalls

2. More hardware techniques to improve throughput

3. Software techniques to improve throughput

Revisiting hazards

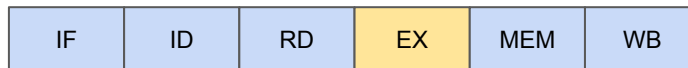When architects became less conservative

Revisiting data hazards: Types of dependencies

- Read After Write (RAW)

  ```
  add R1, R1, #1
  ```

IF	ID	RD	EX	MEM	WB

  ```
  add R3, R1, R2
  ```

	IF	ID	RD	EX	MEM	WB

- Write After Read (WAR)

  ```
  add R3, R1, R2
  ```

IF	ID	RD	EX	MEM	WB

  ```
  load R1, (R2)
  ```

	IF	ID	RD	EX	MEM	WB

- Write After Write (WAW)

  ```
  add R1, R1, #1
  ```

IF	ID	RD	EX	MEM	WB

  ```
  load R1, (R2)
  ```

 | | IF | ID | RD | EX | MEM | WB |
  |----|----|----|----|----|-----|-----|

Revisiting data hazards: Bypassing

- Result of first instruction is produced at the output of stage EX
- Source of second instruction is required at the input of stage EX

- Limiting factor: source operands can only be read from register file (after stage WB)
- Workaround: add a path that forwards/bypasses data from producer to consumer

- Without bypass → Stall three cycles

add R1, R1, #1

IF	ID	RD	EX	MEM	WB

add R3, R1, R2

	IF	ID	ID	ID	ID	RD	EX	MEM	WB

- With bypass →No stalls

add R1, R1, #1

IF	ID	RD	EX	MEM	WB

add R3, R1, R2

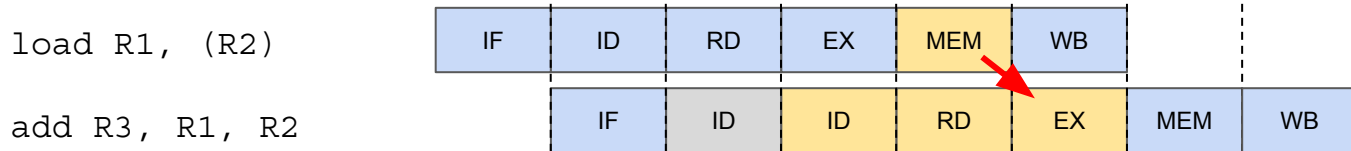|	IF	ID	RD	EX	MEM	WB

6

Revisiting data hazards: Bypassing (cont.)

- Bypassing helps reducing pipeline stalls due to data hazards
- Not all stalls can be avoided

- Without bypass → Stall three cycles

load R1, (R2)									

- With bypass → Stall one cycle
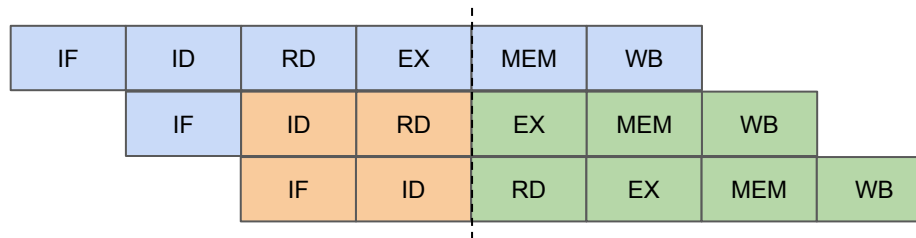
Revisiting control hazards: Branch prediction

■ Instead of stalling the pipeline, we can assume that branches are never taken

- Prediction correct → No stalls

```
beq R1, myfunc
add R3, R1, R2
mul R3, R3, #2
```
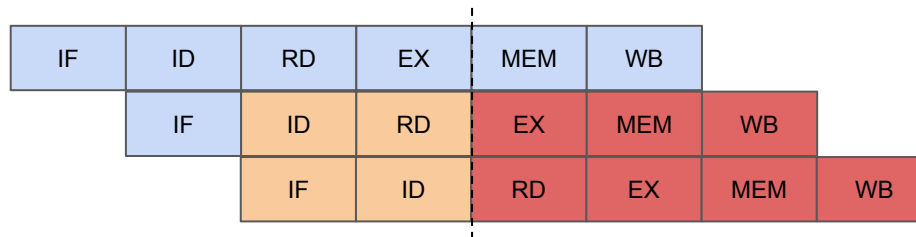
IF	ID	RD	EX	MEM	WB			
	IF	ID	RD	EX	MEM	WB		
		IF	ID	RD	EX	MEM	WB	

- Prediction incorrect → Need to kill the instructions in-flight!!

```
beq R1, myfunc
add R3, R1, R2
mul R3, R3, #2
```

| IF | ID | RD | EX | MEM | WB | | | |
| | IF | ID | RD | EX | MEM | WB | | |
| | | IF | ID | RD | EX | MEM | WB | |

Revisiting control hazards: Branch prediction (cont.)

- There are two aspects which need to be predicted
  - If branch is taken or not
  - What is the target instruction of the branch
- Branch prediction is a very active topic in computer architecture research

- For this course you need to know
  - Branch prediction exists and has an impact on instruction throughput
  - Execution based on branch prediction is called _Speculative Execution_
  - Speculative execution requires a mechanism to kill in-flight instructions

- If you want to learn more
  - Chapter 3.3 of Computer Architecture: A Quantitative Approach
  - Look into the proceedings of almost any computer architecture conference

Revisiting structural hazards: Adding resources

- Example: Reading two operands from the register file during the same cycle
  - Limitation: My register file only has one read port
  - Current implementation: Spend two cycles reading operands

IF	ID	RD1	RD2	EX	MEM	WB

  - Workaround: Add another read port
  - Benefit: Only one cycle to read operands
  - Drawback: More expensive design

| IF	ID	RD	EX	MEM	WB

Multi-cycle pipelines

When architects decided diversify

Different stages for different instructions

- Until now, all instructions follow the same pipeline
- Arithmetic instructions go through the MEM stage without it being necessary
- Proposal: Split pipeline into multiple paths and choose one path after decoding
- Also allows to separate costly operations (eg. multiplications) into their own pipeline

| IF | ID |

| AGU | MEM | WB | `load, store`

| ALU | WB | `add, sub`

| MUL1 | MUL2 | MUL3 | WB | `mul`
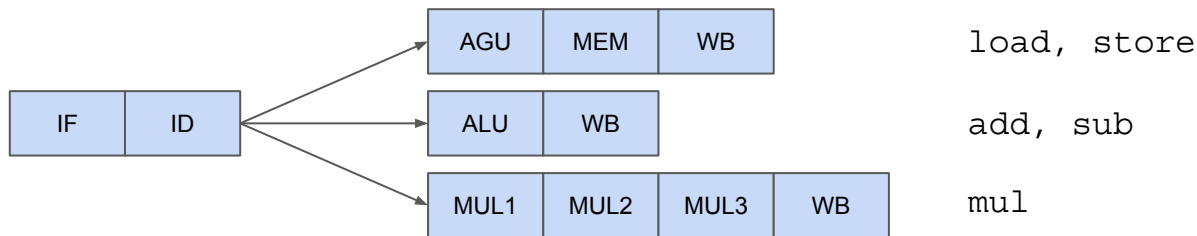
- Benefit: Reducing latency of some operations
- Disclaimer: We are still issuing one instruction per cycle and preserving program order

Dynamic Scheduling

When architects decided to play God
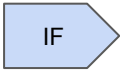
Revisiting data hazards: Beyond bypassing

- Not all data hazards can be circumvented with bypassing
- Pipeline stalls may block later instructions that do not have a dependency

```
1. load R1, (R2)
2. add  R3, R1, R2
3. add  R2, #1
```

| IF	ID	RD	EX	MEM	WB			
	IF	ID	ID	RD	EX	MEM	WB	
		IF	IF	ID	RD	EX	MEM	WB

- There is a data dependency between 1 and 2
- There is NO dependency between 3 and any other previous instruction

- Why should 3 be stalled??
- Proposal: Issue 3 while 2 is waiting for the dependency to be resolved

14

Out-of-Order execution

- Deciding to alter the order of instruction execution is called *Dynamic Scheduling*
- Another common word for this technique is *Out-of-Order* execution (OoO)

- OoO allows us to circumvent data dependencies that cause data hazards
- OoO does NOT circumvent other hazards (eg. structural hazards)

- Fetched instructions are placed into a queue or buffer waiting to be picked by the dynamic scheduler   IF

- When an instruction finishes execution (result available), it has *completed*
- When the instruction result is written and cannot be undone, it has been *retired*

Out-of-Order execution: New headaches

- Upon decoding 2, a data hazard is discovered and the instruction is stalled
- Next cycle, 3 is decoded and issued since it does not have a data dependency

```
1.  mul    R1, R1, #2
2.  add    R3, R1, R2
3.  load   R2, (R4)
```

IF		ID	RD	MUL1	MUL2	MUL3	WB	
	IF		ID			ID	RD	EX	WB
		IF		ID	RD	AGU	MEM	WB

- Sadly, 2 and 3 try to WB during the same cycle → Structural hazard
- Correct OoO execution must stall 2 an extra cycle to avoid the hazard

```
1.  mul    R1, R1, #2
2.  add    R3, R1, R2
3.  load   R2, (R4)
```

IF		ID	RD	MUL1	MUL2	MUL3	WB		
	IF		ID		ID	ID	RD	EX	WB
		IF		ID	RD	AGU	MEM	WB

Out-of-Order execution: Revisiting dependencies

- In an in-order pipeline, the only data dependency that could cause a hazard was a _true dependence_ (RAW)

- OoO execution introduces the possibility of
    - _Anti-dependence_ (WAR)
    - _Output dependence_ (WAW)

New issues → New solutions: Register Renaming

- Proposal: Implement more physical registers than the ISA defines
  - Logical registers (ISA-defined): `R0 - R31`
  - Physical registers (implementation-defined): `P0 - P63`

- The decoding logic renames the logical registers of all instructions to physical registers
- Each write to a logical register produces a new physical register translation

```
mul    R1, R1, #2                    mul    P1, P0, #2
...                  ⟹              ...
add    R1, R1, R3                    add    P2, P1, P3
```

- Register renaming eliminates the WAR and WAW hazards

- If you want to know more, read about *Tomasulo's Algorithm*

Out-of-Order + Speculative execution = 🦠☠️💔

- In-order speculative execution introduced the need for a killing mechanism

- With OoO, speculative instructions can complete before the prediction is even resolved
- How can we undo writes to registers and/or memory?
- What about exceptions? (completely outside the scope of this course)

- There is now easy answer…
  - For registers, register renaming helps
  - For memory, some approaches implement Load and Store Buffers (LB, SB) that hold petitions from/to memory until the prediction is resolved

Superscalar Pipelines

Beyond IPC = 1

Revisiting multi-cycle pipelines

- Having implemented:
  - Multi-cycle pipeline (different pipelines for different types of instructions)
  - Out-of-Order execution
- Why stop at one issue per cycle?
- If no hazards → issue logic can feed into one or more pipelines on a single cycle

- When issuing more than one instruction, we say that it is a *Superscalar Pipeline*
- The benefit is limited if the pipeline can only retire one instruction per cycle

- Modern processors are able to *issue* and *retire* more than one instruction per cycle

Real use case: Intel Skylake

- **Front-End**
  - Out-of-Order execution
  - Multiple fetch per cycle
  - Register renaming

- **Multiple issue per cycle**

- **Back-End**
  - Eight execution pipelines
  - Load Buffer with 72 entries
  - Store Buffer with 56 entries
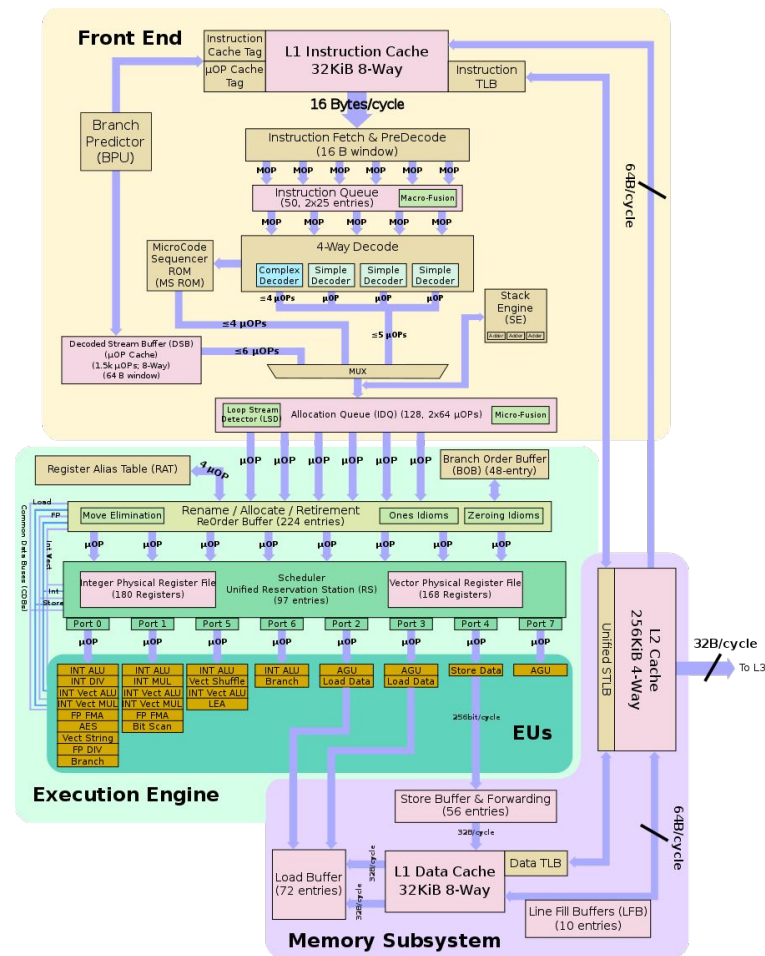
Software techniques to improve ILP

When programmers try to help architects

Basic block

Straight-line code sequence with no branches in except to the entry and no branches out except at the exit.

Chapter 3 of Computer Architecture: A Quantitative Approach

- Sequence of instructions that are likely to depend upon each other
- Small blocks with lots of dependencies limit the overlap between instructions
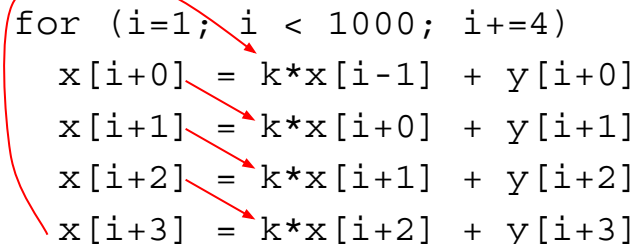
- Example: Body of a loop

```
for (i=0; i < 1000; i++)
   x[i] = x[i] + y[i]
```

```
body:
   load  R2, @x[i]
   load  R3, @y[i]
   add   R2, R2, R3
   store @x[i], R2
   add   R1, 1
   beq   R1, body
```

1. Increase size of basic block
2. Leverage Instruction Level Parallelism (ILP) across multiple basic blocks

Loop unrolling

- Technique to increase the basic block size
  - Benefit: Exposes more ILP
  - Drawbacks: Less instruction cache locality, increased register pressure
- Limitations
  - If there are dependencies across iterations, there is little to be gained
  - Number of iterations must be multiple of the unroll factor

```
for (i=0; i < 1000; i+=4)
  x[i+0] = k*x[i+0] + y[i+0]
  x[i+1] = l*x[i+1] + y[i+1]
  x[i+2] = k*x[i+2] + y[i+2]
  x[i+3] = k*x[i+3] + y[i+3]
```

```
for (i=1; i < 1000; i+=4)
  x[i+0] = k*x[i-1] + y[i+0]
  x[i+1] = k*x[i+0] + y[i+1]
  x[i+2] = k*x[i+1] + y[i+2]
  x[i+3] = k*x[i+2] + y[i+3]
```

Software pipelining

- Reorganize loops such that each iteration contains instructions from different iterations of the original code.
- Very tricky to implement → Let the compiler handle it

(a) Software pipelining

(b) Loop unrolling

```
load   R1, @k


load   R2, @x[i+0]
load   R3, @x[i+0]
mul    R4, R1, R2
add    R5, R4, R3
store @x[i+0], R5

load   R6, @x[i+1]
load   R7, @x[i+1]
mul    R8, R1, R6
add    R9, R8, R7
store @x[i+0], R9

...
```

Remarks about software techniques

1. Simplify code structure
   - Minimize conditional and nested regions of code
   - Be aware of (and avoid) unnecessary data dependencies

2. Don't overdo it
   - Some optimizations (eg. software pipelining) can be done by the compiler

3. Hardware implementation is smarter than you think
   - Dynamic scheduling is able to do runtime re-ordering, which is not possible when writing code (not even from the compiler's point of view)
   - There are many other hardware optimization techniques that we have not covered today!