

Instruction Level Parallelism

Tricking the system to maximize throughput

Source material

- D. A. Patterson and J. L. Hennessy, *Computer Organization and Design: The Hardware/Software Interface: RISC-V Edition*. Cambridge, MA: Morgan Kaufmann, 2018.
 - Chapter 4.7: Data Hazards: Forwarding versus Stalling
 - Chapter 4.8: Control Hazards
 - Chapter 4.10: Parallelism via Instructions
- J. L. Hennessy, D. A. Patterson, and K. Asanović, *Computer Architecture: A Quantitative Approach*, 5th ed. Waltham, MA: Morgan Kaufmann/Elsevier, 2012.
 - Appendix C: Pipelining: Basic and Intermediate Concepts
 - Chapter 3: Instruction-Level Parallelism and Its Exploitation

Goal: Improving throughput by increasing ILP

1. Revisit how we circumvent hazards to avoid pipeline stalls
2. More hardware techniques to improve throughput
3. Software techniques to improve throughput

Revisiting hazards

When architects became less conservative

Revisiting data hazards: Types of dependencies

- Read After Write (RAW)

add R1, R1, #1



add R3, R1, R2



- Write After Read (WAR)

add R3, R1, R2



load R1, (R2)



- Write After Write (WAW)

add R1, R1, #1




load R1, (R2)




Revisiting data hazards: Bypassing

- Result of first instruction is produced at the output of stage EX
 - Source of second instruction is required at the input of stage EX
 - Limiting factor: source operands can only be read from register file (after stage WB)
 - Workaround: add a path that forwards/bypasses data from producer to consumer
- Without bypass → Stall three cycles



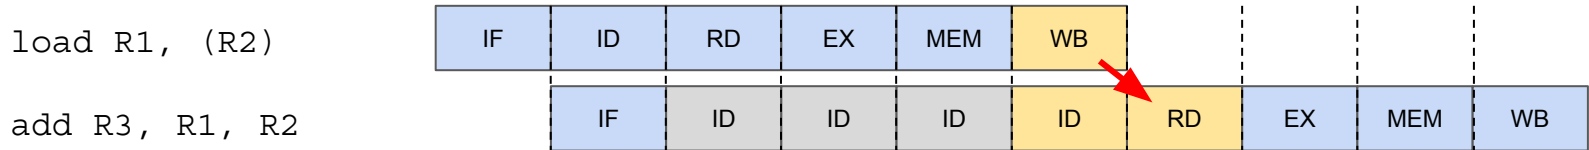
- With bypass → No stalls



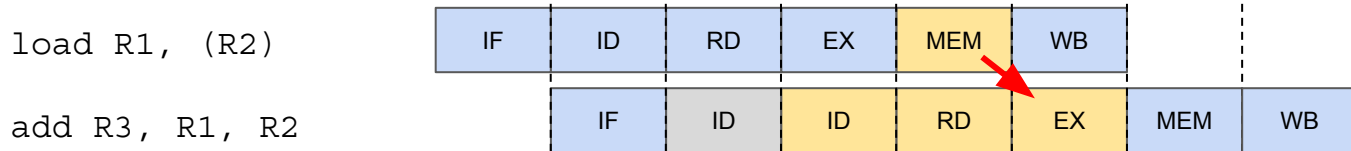
Revisiting data hazards: Bypassing (cont.)

- Bypassing helps reducing pipeline stalls due to data hazards
- Not all stalls can be avoided

- Without bypass → Stall three cycles



- With bypass → Stall one cycle



Revisiting control hazards: Branch prediction


- Instead of stalling the pipeline, we can assume that branches are never taken

- Prediction correct → No stalls

beq R1, myfunc

add R3, R1, R2

mul R3, R3, #2




- Prediction incorrect → Need to kill the instructions in-flight!!

beq R1, myfunc

add R3, R1, R2

mul R3, R3, #2

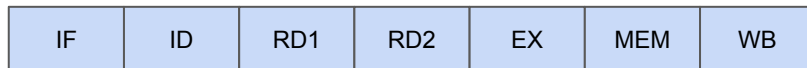


Revisiting control hazards: Branch prediction (cont.)

- There are two aspects which need to be predicted
 - If branch is taken or not
 - What is the target instruction of the branch
- Branch prediction is a very active topic in computer architecture research
- For this course you need to know
 - Branch prediction exists and has an impact on instruction throughput
 - Execution based on branch prediction is called Speculative Execution
 - Speculative execution requires a mechanism to kill in-flight instructions
- If you want to learn more
 - Chapter 3.3 of Computer Architecture: A Quantitative Approach
 - Look into the proceedings of almost any computer architecture conference

Revisiting structural hazards: Adding resources

- Example: Reading two operands from the register file during the same cycle
 - Limitation: My register file only has one read port
 - Current implementation: Spend two cycles reading operands



- Workaround: Add another read port
- Benefit: Only one cycle to read operands
- Drawback: More expensive design

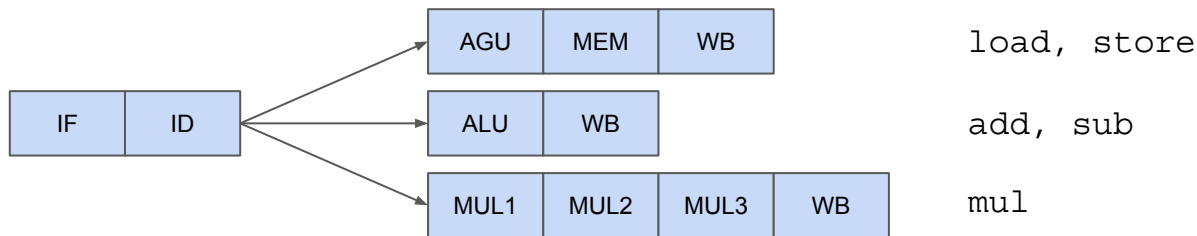


Multi-cycle pipelines

When architects decided to diversify

Different stages for different instructions

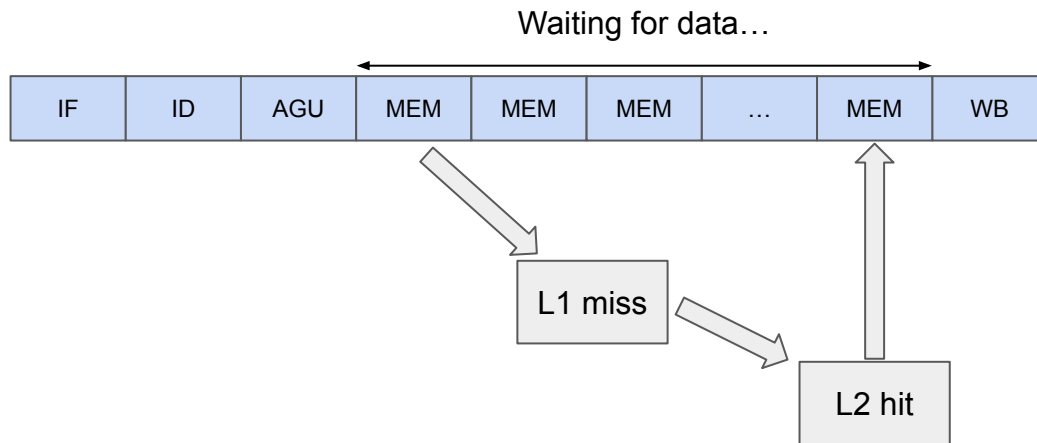
- Until now, all instructions follow the same pipeline
- Arithmetic instructions go through the MEM stage without it being necessary
- Proposal: Split pipeline into multiple paths and choose one path after decoding
- Also allows to separate costly operations (eg. multiplications) into their own pipeline



- Benefit: Reducing latency of some operations
- Disclaimer: We are still issuing one instruction per cycle and preserving program order

Number of cycles is not always known at decode

- In micro-architectures that include caches, how many cycles will a load cost?
 - L1 Hit → 1 cycle? 3 cycles?
 - L1 Miss, L2 Hit → 20 cycles?
 - Going to main memory → 100 cycles?
- Special logic has to be added to the MEM stage to wait for load data to arrive



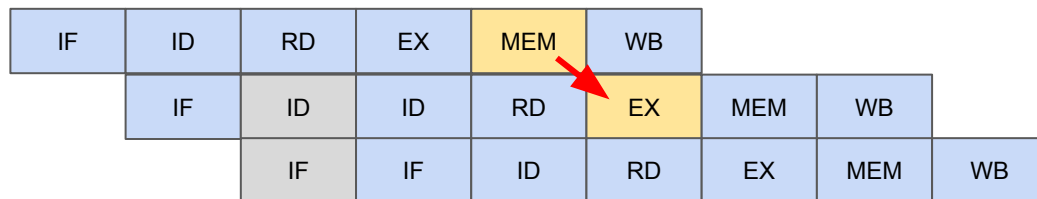
Dynamic Scheduling

When architects decided to play God

Revisiting data hazards: Beyond bypassing

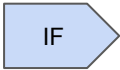
- Not all data hazards can be circumvented with bypassing
- Pipeline stalls may block later instructions that do not have a dependency

```
1. load R1, (R2)
2. add  R3, R1, R2
3. add  R2, #1
```




- There is a data dependency between 1 and 2
- There is NO dependency between 3 and any other previous instruction
- Why should 3 be stalled??
- Proposal: Issue 3 while 2 is waiting for the dependency to be resolved

Out-of-Order execution


- Deciding to alter the order of instruction execution is called Dynamic Scheduling
- Another common word for this technique is Out-of-Order execution (OoO)
- OoO allows us to circumvent data dependencies that cause data hazards
- OoO does NOT circumvent other hazards (eg. structural hazards)
- Fetched instructions are placed into a queue or buffer waiting to be picked by the dynamic scheduler 
- When an instruction finishes execution (result available), it has completed
- When the instruction result is written and cannot be undone, it has been retired

Out-of-Order execution: New headaches

- Upon decoding 2, a data hazard is discovered and the instruction is stalled
- Next cycle, 3 is decoded and issued since it does not have a data dependency



- Sadly, 2 and 3 try to WB during the same cycle → Structural hazard
- Correct OoO execution must stall 2 an extra cycle to avoid the hazard



Out-of-Order execution: Revisiting dependencies

- In an in-order pipeline, the only data dependency that could cause a hazard was a true dependence (RAW)
- OoO execution introduces the possibility of
 - Anti-dependence (WAR)
 - Output dependence (WAW)

New issues → New solutions: Register Renaming

- Proposal: Implement more physical registers than the ISA defines
 - Logical registers (ISA-defined): R0 - R31
 - Physical registers (implementation-defined): P0 - P63
- The decoding logic renames the logical registers of all instructions to physical registers
- Each write to a logical register produces a new physical register translation

mul	R1	R1	#2		mul	P1,	P0	#2
...				→	...			
add	R1	R1	R3		add	P2	P1	P3

- Register renaming eliminates the WAR and WAW hazards
- If you want to know more, read about *Tomasulo's Algorithm*

Out-of-Order + Speculative execution =

- In-order speculative execution introduced the need for a killing mechanism
- With OoO, speculative instructions can complete before the prediction is even resolved
- How can we undo writes to registers and/or memory?
- What about exceptions? (completely outside the scope of this course)
- There is now easy answer...
 - For registers, register renaming helps
 - For memory, some approaches implement Load and Store Buffers (LB, SB) that hold petitions from/to memory until the prediction is resolved

Superscalar Pipelines

Beyond $IPC = 1$

Revisiting multi-cycle pipelines

- Having implemented:
 - Multi-cycle pipeline (different pipelines for different types of instructions)
 - Out-of-Order execution
- Why stop at one issue per cycle?
- If no hazards → issue logic can feed into one or more pipelines on a single cycle
- When issuing more than one instruction, we say that it is a Superscalar Pipeline
- The benefit is limited if the pipeline can only retire one instruction per cycle
- Modern processors are able to issue and retire more than one instruction per cycle

Real use case: Intel Skylake

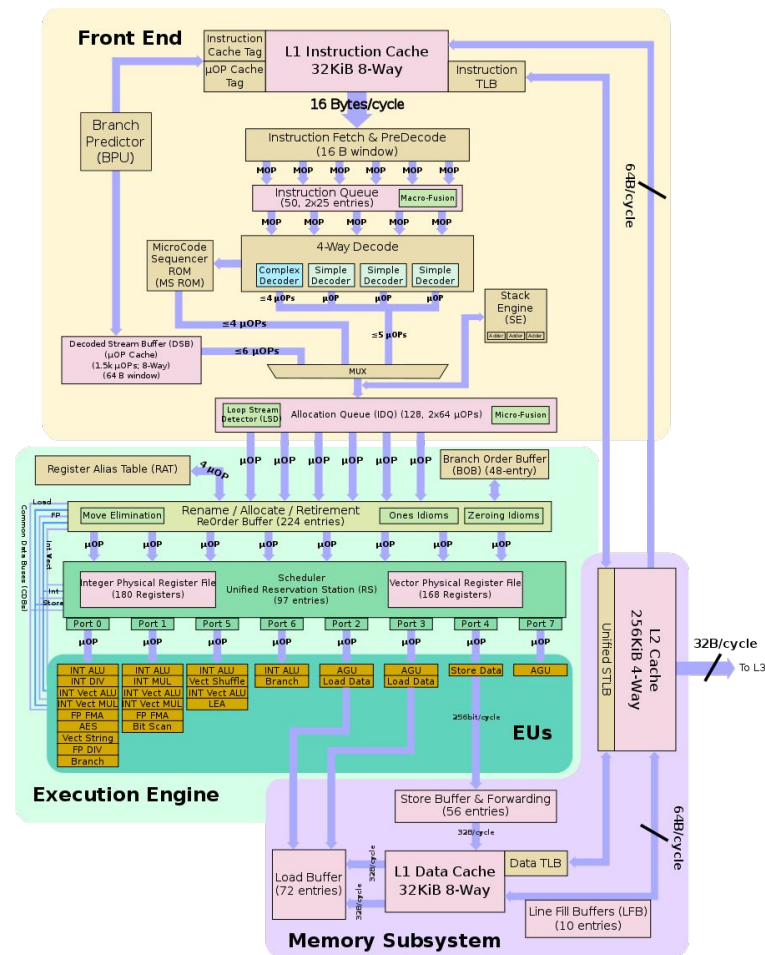
■ Front-End

- Out-of-Order execution
- Multiple fetch per cycle
- Register renaming

■ Multiple issue per cycle

■ Back-End

- Eight execution pipelines
- Load Buffer with 72 entries
- Store Buffer with 56 entries



Calculating theoretical FP peak

■ Arithmetic (eg. Floating-Point) performance

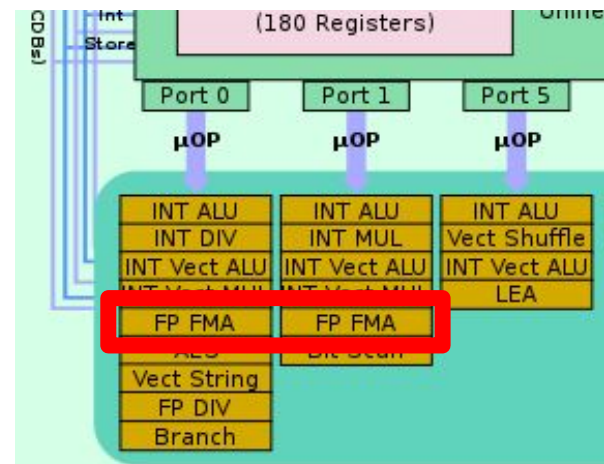
- P := Number of Pipelines
- U := Number of Functional units (for SIMD, width of SIMD registers)
- T := Throughput of instruction in ONE pipeline
- O := Operations per instruction
- F := CPU frequency

■ Real use case: Intel Skylake

- P = 2 FMA pipelines
- U = 8 DP elements per SIMD register (AVX512)
- T = 1 FMA per cycle
- O = 2 operations per instruction
- F = 2.10 GHz (cycles / second)

■ Peak performance



- ~~2 * 8 elements/instruction * 1 instruction/cycle * 2 ops/element~~ = 32 ops/cycle
- 32 ops/cycle * 2.10 * 10⁹ cycles/second = 67.20 * 10⁹ ops/s = **67.20 GFlop/s**





Another example: Cortex-A57 (Jetson TX1)

- <https://chipsandcheese.com/p/cortex-a57-nintendo-switchs-cpu>

- DP Scalar instructions (aarch64)



- DP SIMD instructions (NEON)



Software techniques to improve ILP

When programmers try to help architects

Basic block

Straight-line code sequence with no branches in except to the entry and no branches out except at the exit.

Chapter 3 of Computer Architecture: A Quantitative Approach

- Sequence of instructions that are likely to depend upon each other
- Small blocks with lots of dependencies limit the overlap between instructions

- Example: Body of a loop

```
for (i=0; i < 1000; i++)  
    x[i] = x[i] + y[i]
```

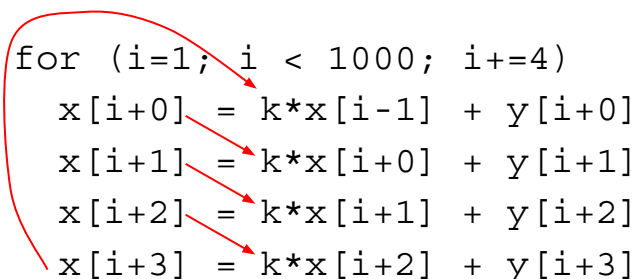
```
body:  
    load    R2, @x[i]  
    load    R3, @y[i]  
    add     R2, R2, R3  
    store   @x[i], R2  
    add     R1, 1  
    beq     R1, body
```

1. Increase size of basic block
2. Leverage Instruction Level Parallelism (ILP) across multiple basic blocks

Loop unrolling

- Technique to increase the basic block size
 - Benefit: Exposes more ILP
 - Drawbacks: Less instruction cache locality, increased register pressure
- Limitations
 - If there are dependencies across iterations, there is little to be gained
 - Number of iterations must be multiple of the unroll factor

```
for (i=0; i < 1000; i+=4)
    x[i+0] = k*x[i+0] + y[i+0]
    x[i+1] = l*x[i+1] + y[i+1]
    x[i+2] = k*x[i+2] + y[i+2]
    x[i+3] = k*x[i+3] + y[i+3]
```

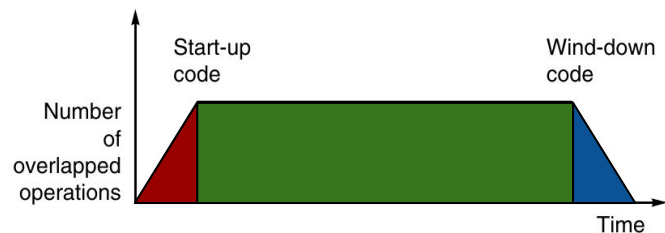


```
for (i=1; i < 1000; i+=4)
    x[i+0] = k*x[i-1] + y[i+0]
    x[i+1] = k*x[i+0] + y[i+1]
    x[i+2] = k*x[i+1] + y[i+2]
    x[i+3] = k*x[i+2] + y[i+3]
```

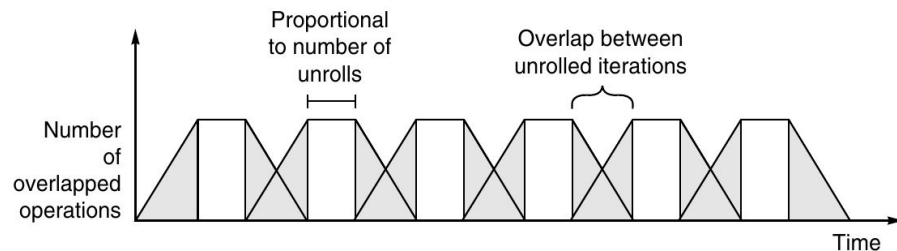
The diagram illustrates the unrolled loop with a red circle around the first iteration and red arrows pointing from the right-hand side of one iteration to the left-hand side of the next, highlighting the data dependencies between iterations.

Software pipelining

- Reorganize loops such that each iteration contains instructions from different iterations of the original code.
- Very tricky to implement → Let the compiler handle it



(a) Software pipelining



(b) Loop unrolling

```
load  R1, @k
```

```
load  R2, @x[i+0]
```

```
load  R3, @x[i+0]
```

```
mul   R4, R1, R2
```

```
add   R5, R4, R3
```

```
store @x[i+0], R5
```

```
load  R6, @x[i+1]
```

```
load  R7, @x[i+1]
```

```
mul   R8, R1, R6
```

```
add   R9, R8, R7
```

```
store @x[i+0], R9
```

...

Remarks about software techniques

1. Simplify code structure

- Minimize conditional and nested regions of code
- Be aware of (and avoid) unnecessary data dependencies

2. Don't overdo it

- Some optimizations (eg. software pipelining) can be done by the compiler

3. Hardware implementation is smarter than you think

- Dynamic scheduling is able to do runtime re-ordering, which is not possible when writing code (not even from the compiler's point of view)
- There are many other hardware optimization techniques that we have not covered today!