

# **Distributed Memory**

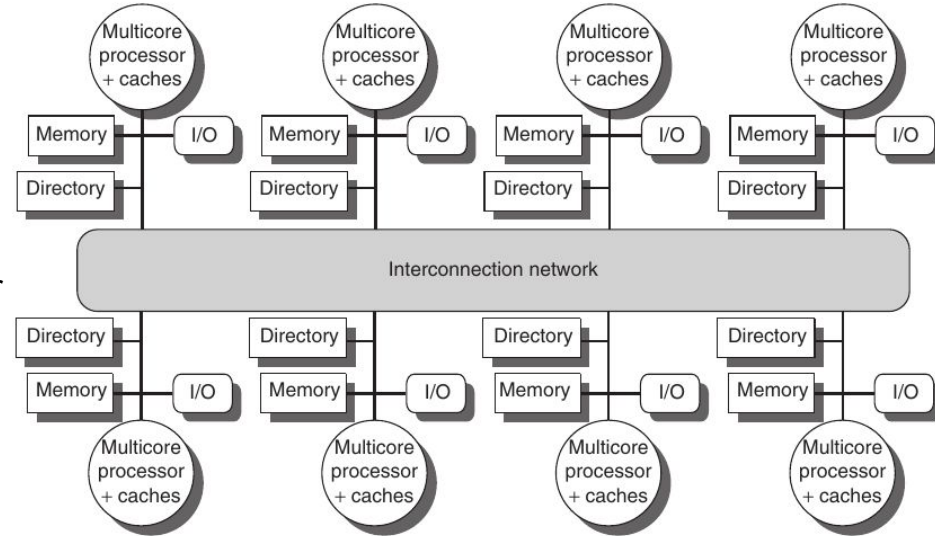
Too many people in the same room

# Source material

- J. L. Hennessy, D. A. Patterson, and K. Asanović, *Computer Architecture: A Quantitative Approach*, 5th ed. Waltham, MA: Morgan Kaufmann/Elsevier, 2012.
  - Chapter 6: Warehouse-Scale Computers to Exploit Request-Level and Data-Level Parallelism
  - Appendix F: Interconnection Networks
- D. A. Patterson and J. L. Hennessy, *Computer Organization and Design: The Hardware/Software Interface: RISC-V Edition*. Cambridge, MA: Morgan Kaufmann, 2018.
  - Chapter 6: Parallel Processors from Client to Cloud
- [R. Copetti, \*Architecture of Consoles - A practical analysis\*](#)

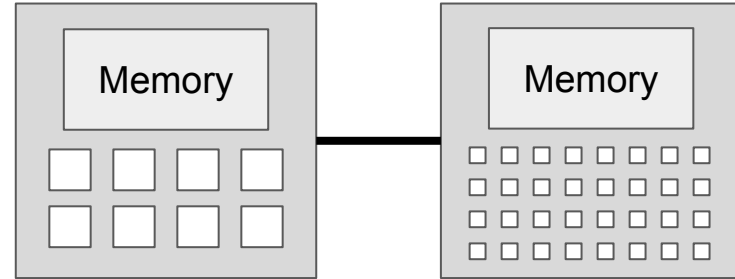
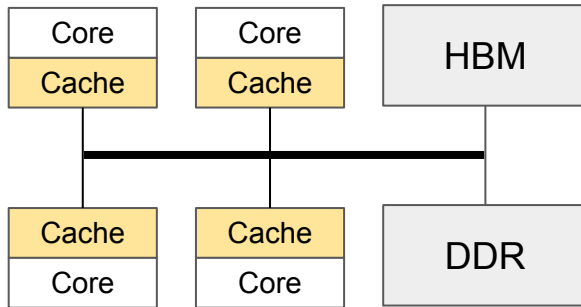
# Non-Uniform Memory Access (NUMA)

- Number of processors in a system increase
  - More petitions to/from memory
  - More network contention
  - Less bandwidth per processor
- Break-up memory into slices
  - Each slices is positioned close to a processor
  - Access symmetry is broken
  - Hence, Non-Uniform Memory Access
- Threads can still share the memory space
- OS knows NUMA topology and provides mechanisms to work with it
  - When software is capable of distinguishing NUMA regions, we call it NUMA-aware



# Heterogeneous memories

- Complex systems with different types of memory technologies
  - Eg: DDR + HBM
  - Eg: Host - Device
- OS provides software layer that can be
  - Implicit (or transparent): Software is not aware of which physical memory is using
  - Explicit: Software must specify which physical memory wants to access



# Memories from a distant land

- As systems become bigger and processors have to access farther memories, the cost of maintaining consistency and coherence becomes too high
- Hence, the address space is broken into isolated partitions
  - Threads from the same process still share memory address space
  - Processes, which are managed by the OS, do not share memory address space
- This is what we commonly refer to as Distributed Memory
- New challenges: How do processes share data?
  - A common way to do so is with Message Passing
- New possibilities: Programs are not bound by the memory capacity of one node
  - Programs be multi-process and leverage memory of multiple nodes

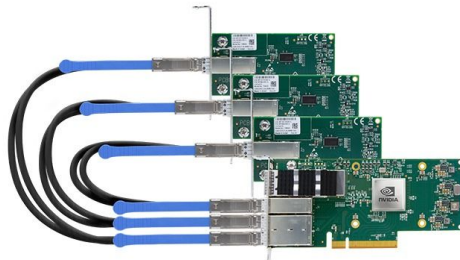
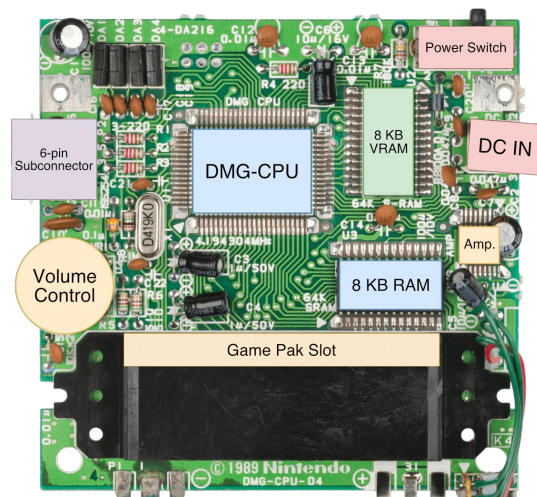
# Hardware support to move data around

- Moving data between locations in a distributed memory system takes time
- Naive approach: have the processor read the data from its memory and send it through the network to the remote memory
  - If the amount of data is very big it will require multiple transactions
  - Even if data is small, the transfers could be very frequent
  - Result: the processor spends most of its cycles moving data from memory locations
- Workaround: implement a dedicated hardware component to move data around
  - This component is called Direct Memory Access (DMA) engine
  - Give it a source, the size of the transfer, a destination, and kick it into action
  - When the DMA transfer is done, it will notify the processor with an interruption

# Interfacing with a DMA engine

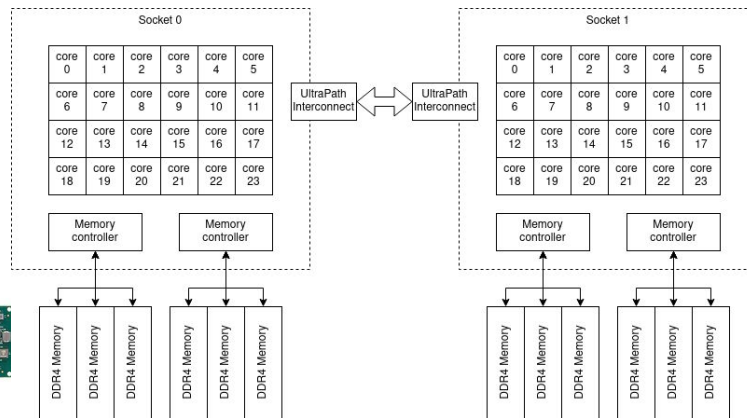
## Embedded systems and low-level code

- Write to memory-mapped registers
- Implement response to DMA interrupt

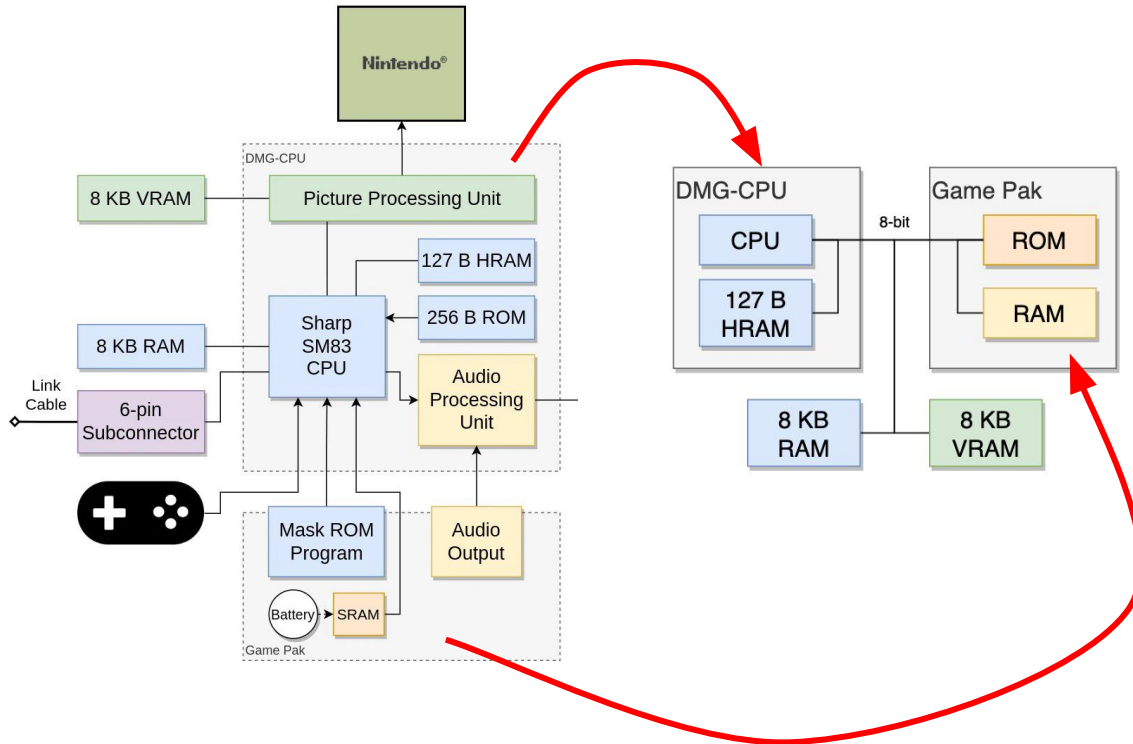


## High-level code

- Completely transparent
  - OS and kernel management (`malloc`)
- Programmer-friendly API
  - `cudamalloc`
  - `MPI_Send`, `MPI_Broadcast`



# Personal favorite example (not HPC)



Interrupt Register	0xFFFF
High RAM	0xFF80 - 0xFFFE
Unusable	0xFF4C - 0xFF7F
I/O	0xFF00 - 0xFF4B
Unusable	0xFE00 - 0xFEFF
Sprite Attributes	0xFE00 - 0xFE9F
Unusable	0xE000 - 0xFDFD
Internal RAM	0xC000 - 0xDFFF
Switchable RAM Bank	0xA000 - 0xBFFF
Video RAM	0x8000 - 0x9FFF
Switchable ROM Bank	0x4000 - 0x7FFF
ROM	0x0000 - 0x3FFF



# Distributed Memory Programming models

- Message Passing Interface (MPI)
  - Very prevalent in CPU-based HPC
  - Explicit message passing with high-level API
- Compute Unified Device Architecture (CUDA)
  - Very common in NVIDIA GPUs
  - Explicit communication with high-level API
- OpenMP Device Offload
  - Extension of the shared-memory programming model
  - Implicit communication with pragmas
- Partitioned Global Address Space (PGAS)
  - Software emulates shared address space across multiple distributed memory nodes
  - Examples: UPC++, Global Arrays

# MPI Programming Model

Calling for help

# Source material

- [MPI: A Message-Passing Interface Standard](#)
- [MPI cheatsheet](#)

# MPI Programming Model

- Distributed-memory programming model
- Process level parallelism
- Explicit process communication
- Programming via MPI calls
- Code is written with MPI execution in mind from the very start

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char* argv[]) {
    int nranks, myrank;

    MPI_Init(&argc, &argv);

    MPI_Comm_size(MPI_COMM_WORLD, &nranks);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);

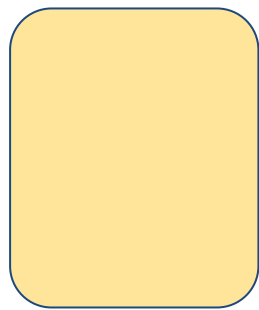
    if (myrank == 0)
    {
        printf("Master rank prints\n");
    }

    MPI_Finalize();
}
```

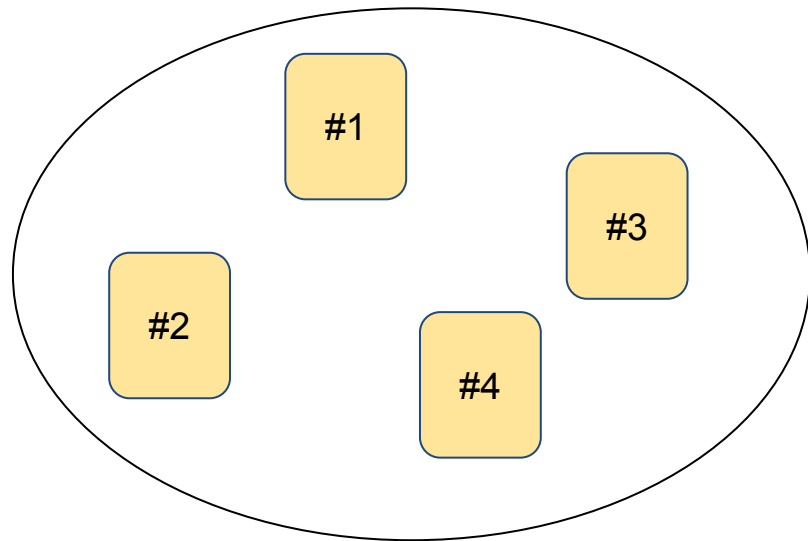
# Vocabulary

## Process $\approx$ Rank

- Program in execution
- Has a Process Identifier (PID)
- Has its own memory space
- Has a rank id
- Groups of ranks are called communicators




Process

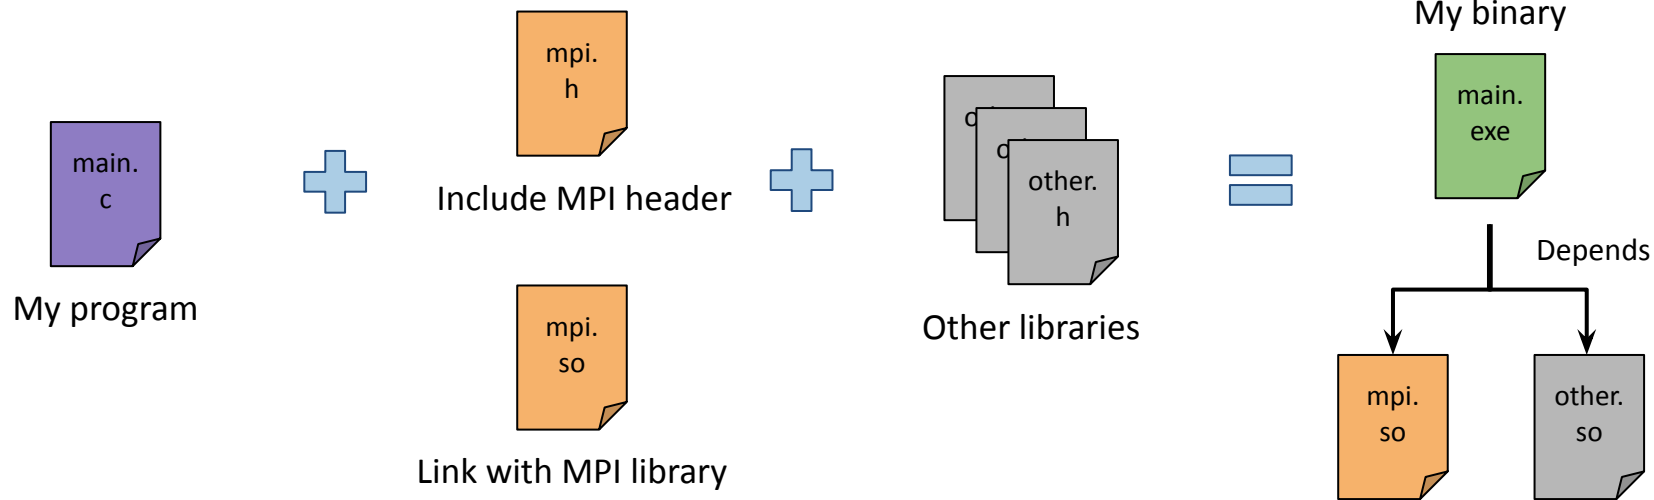


MPI\_COMM\_WORLD communicator

# MPI standards and libraries

- MPI is a standard
  - `MPI_Init`, `MPI_Finalize`
  - `MPI_Barrier`, `MPI_Allreduce`
- There are different library implementations of the same standard
  - OpenMPI, IntelMPI, MPICH
  -  **Disclaimer:** Do not mix the standard OpenMP with the library OpenMPI !!!
- Different libraries implement the same standard with different strategies

# Building an MPI program

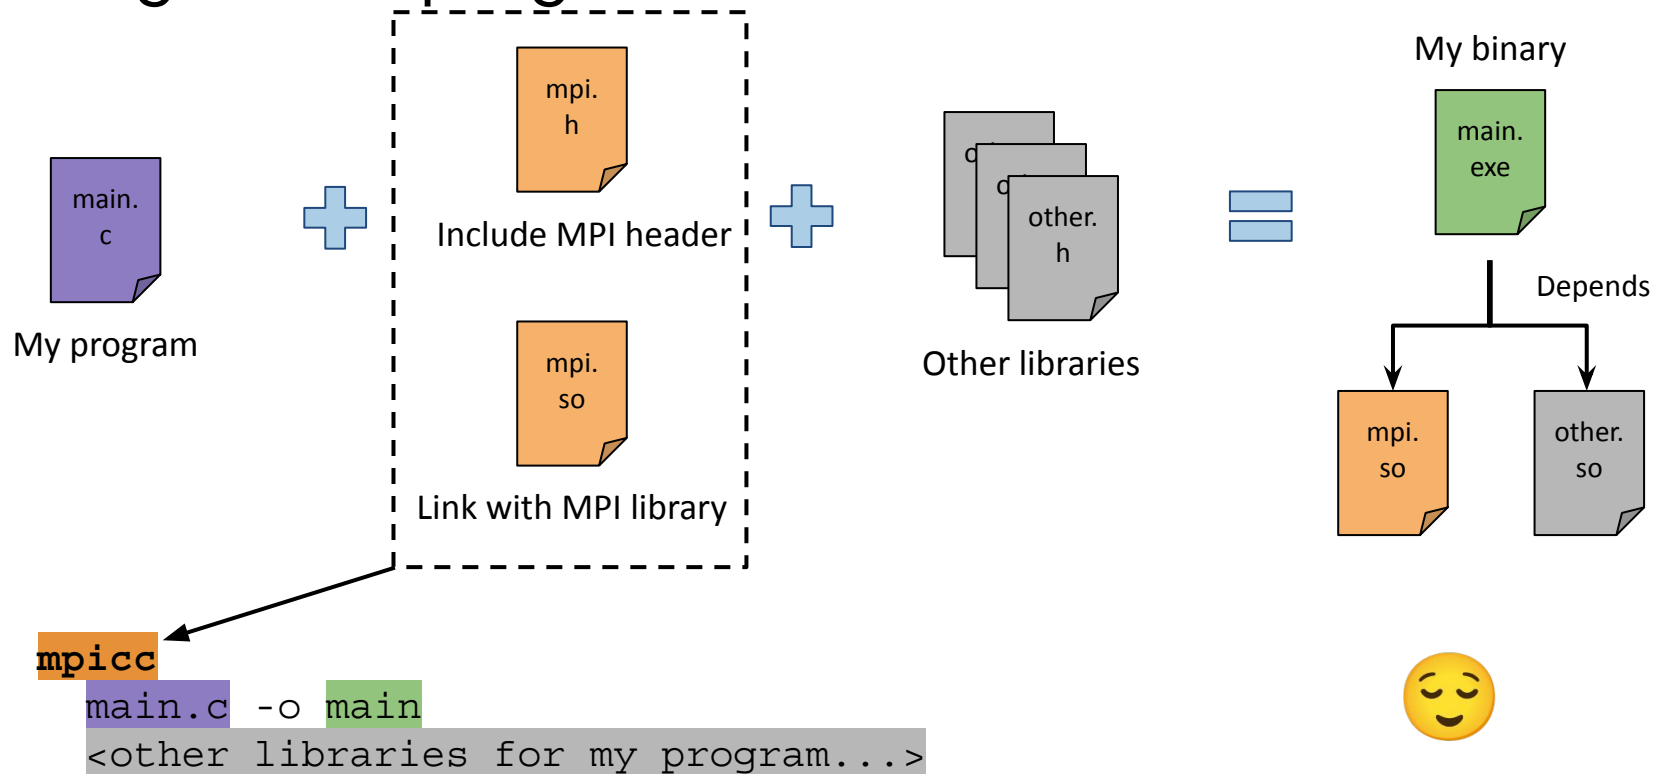


icc

```
main.c -o main  
<other libraries for my program...>  
-I/apps/INTEL/2017.4/impi/2017.3.196/intel64/include  
-L/apps/INTEL/2017.4/impi/2017.3.196/intel64/lib  
-lmpi
```



# Building an MPI program





# MPI parameters and SLURM

- Parameters vary depending on the MPI library
- All MPI libraries have one parameter in common: number of ranks (`-np`)
- Depending on SLURM configuration the number of ranks
  - Will not be set at all (user has to set it)
  - Will be set based on the allocation option `--ntasks`
  - Will be set to a default value (usually the total number of cores)
- Best practices
  - If allocating via SLURM, always set `--ntasks`

# Who participates in an MPI communication

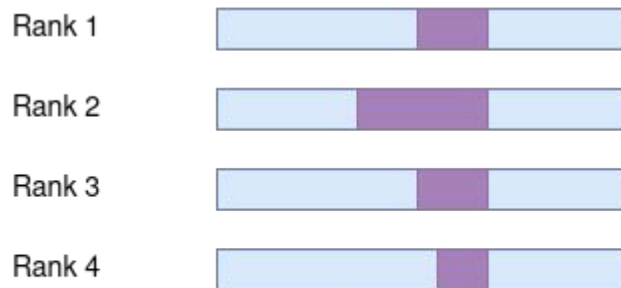
## Point-to-point

- Pair of processes
- `MPI_Send(...)`
- `MPI_Recv(...)`



## Collective

- All processes in a communicator
- `MPI_Barrier(...)`
- `MPI_Allreduce(...)`



Useful computation

MPI\_Send

MPI\_Recv

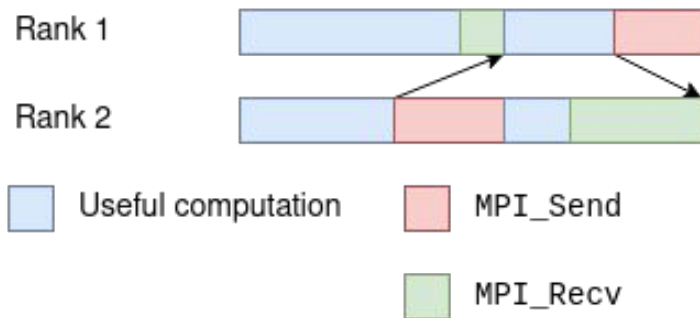
Useful computation

MPI\_Allreduce

# How patient are we during an MPI communication

## Blocking

- Wait until communication has finished
- `MPI_Send(...)`
- `MPI_Recv(...)`



## Non-blocking

- Continue until you really need the data
- `MPI_Isend(...)`
- `MPI_IRecv(...)`
- `MPI_Wait(...)`

