


Memory Hierarchy

Feeding your CPU

The MEM stage


- We have been seeing this diagram of a Pipeline



But where do these memory requests go?
How fast are they?

Simplest design: Main Memory

- CPU's requests go to the external RAM (Random Access Memory)



The main memory is slow because it's:

- Big (hard to index)
- Far (Data must physically travel)


Many instructions depend on memory operations → We need them to be as fast as possible!

... but we can't make the RAM smaller or move it closer...

Are memory accesses “Random”?

- Is using RAM the best option when accesses are not random?

```
int IN[128][8]
int OUT[128][8]
int V[8]
for (int i=0; i<128; ++i){
  for(int j=0; j<8; ++j){
    OUT[i][j] = IN[i][j] * V[j]
  }
}
```



Multiply all rows of “IN” by a vector “V”

Total iterations: $128 \times 8 = 1024$

Accesses **OUT**: 1024 stores

Accesses **IN**: 1024 loads

Accesses **V**: 1024 loads

But V only has 8 elements! We are accessing them over and over...

Temporal Locality

- When something is accessed frequently, we say it has “**Temporal Locality**”
- And what do we do with frequent usage?



VS



VS





VS



We bring them closer, to smaller storage!

- In computer architecture, we call this a **Cache**
- With a small cache, we reduce the memory accesses

```
int IN[128][8]
int OUT[128][8]
int V[8]
for (int i=0; i<128; ++i){
  for(int j=0; j<8; ++j){
    OUT[i][j] = IN[i][j] * V[j]
  }
}
```





After the first 8 accesses to **V**, the CPU gets the data from the cache.

Memory Accesses = 2048 + 8
Cache Accesses = 1016

33% cache usage (hit ratio)...
Can we do better?

Spatial locality

- While `IN[]` and `OUT[]` are **not reused**, they are not random either.
- We are doing consecutive accesses:
- How do we take advantage of this?



Each time we access the RAM, we put N contiguous elements in the Cache.


How big is N ?

- Too big \rightarrow More Silicon, More Risky
- Too small \rightarrow Not useful

Cache Line

- This “N” is called Cache Line, and it’s usually 32, 64, or 128 Bytes wide (*int* = 4 Bytes)

```
int IN[128][8]
int OUT[128][8]
int V [8]
for (int i=0; i<128; ++i){
  for(int j=0; j<8; ++j){
    OUT[i][j] = IN[i][j] * V[j]
  }
}
```



With a Cache Line of 64B, only 1/16 accesses go to the main memory.

IN[] and OUT[]:

Memory accesses: $1/16 \times 2048 = 128$

Cache accesses: $15/16 \times 2048 = 1920$

V[]:

Memory accesses: 1

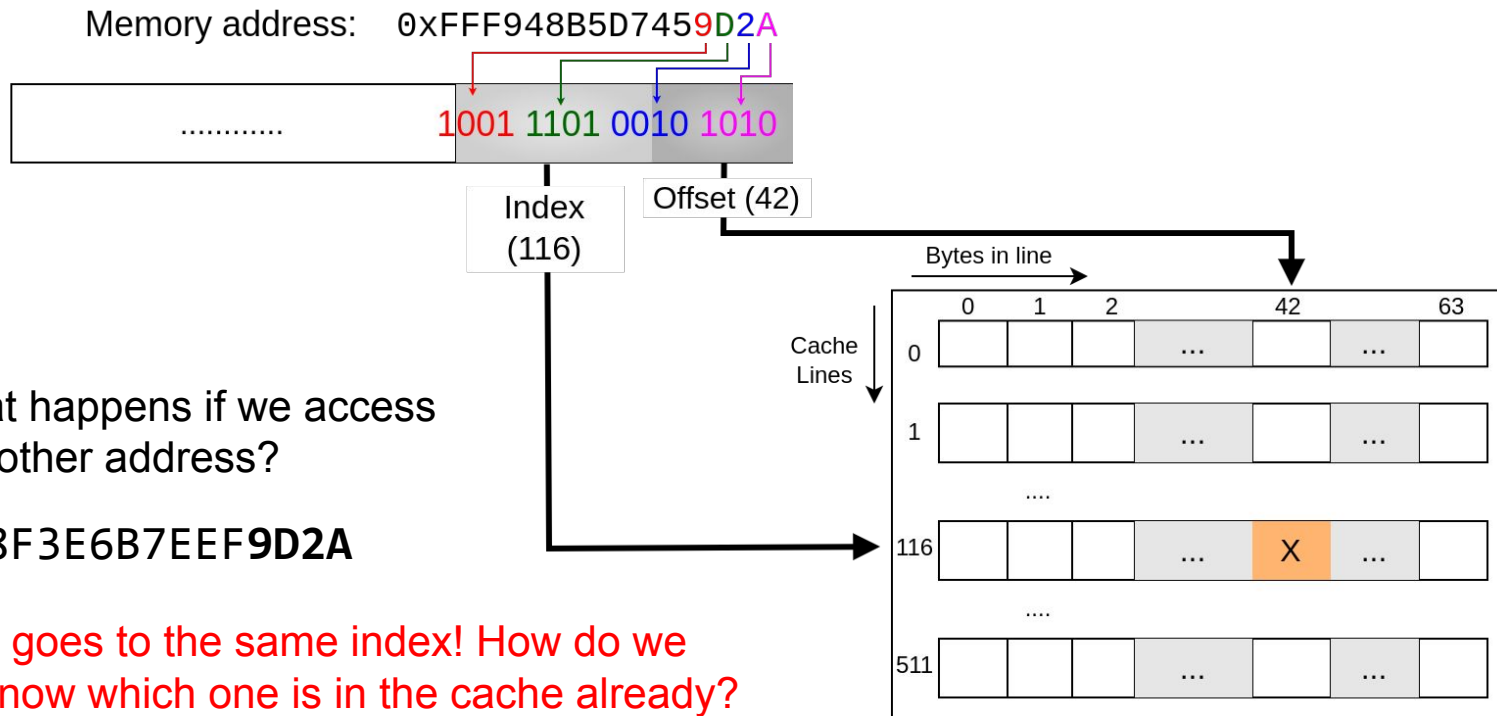
Cache accesses: 1023

Miss ratio:

$$\frac{129}{3072} = 4.2\%$$

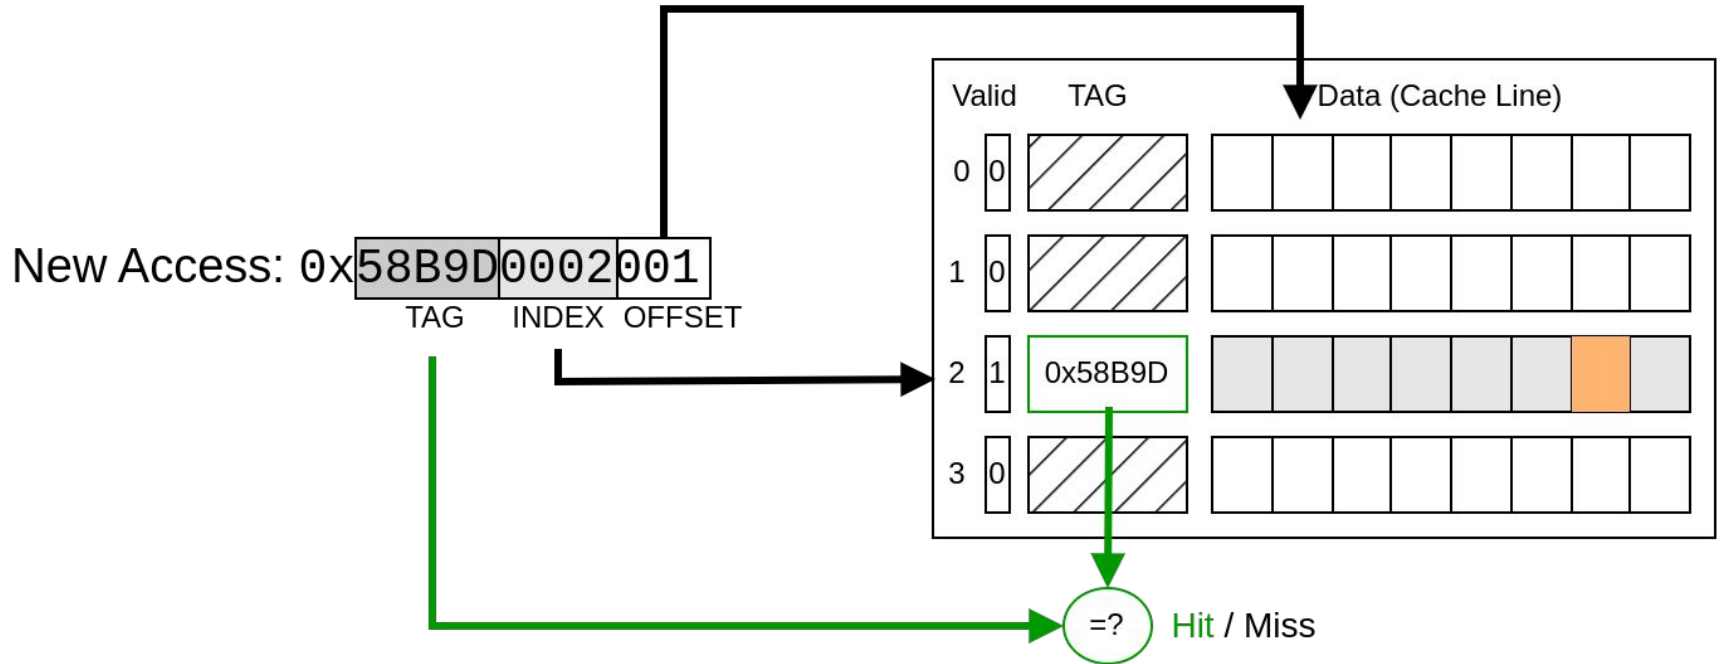
How do we actually fill this cache?

- Easiest approach: **Direct mapping** (lower bits decide where to store data)




Collisions in a Cache

- We need to add a new field to distinguish addresses: the **TAG**



Collisions in a Cache

- We need to add a new field to distinguish addresses: the **TAG**



Verdict on direct mapping

- Benefits:
 - Easy to implement (extracting bits from an address to index a table is fast and cheap)
 - One single comparison per access
 - This allows to have big caches → Bigger problems can fit

- Drawbacks:
 - High collision chance:

`int A[256][32]` → `0xF3B44A57D6010000`
`int B[256][32]` → `0xF3B44A57D6020000` } Consecutive arrays usually collide on lower bits

```
for(int i=0; i<256; ++i){  
    for(int j=0; j<32; ++j){  
        A[i][j] = A[i][j] + B[i][j]  
    }  
}
```

No matter how big the Cache is, each access will evict the other array

↓
All accesses will be misses

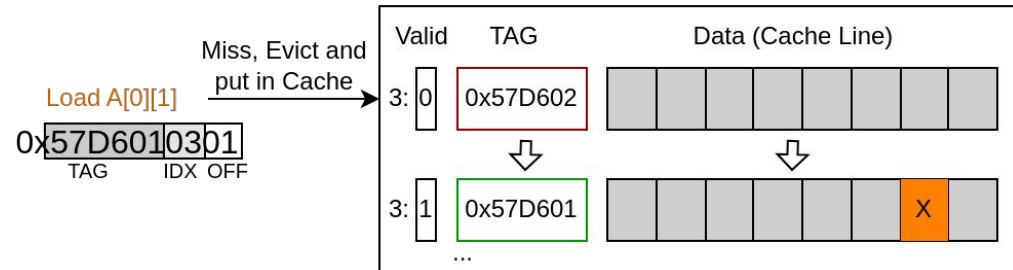

More on Cache Trashing

- A and B share a Cache Line:

```
int A[256][32]
int B[256][32]
```

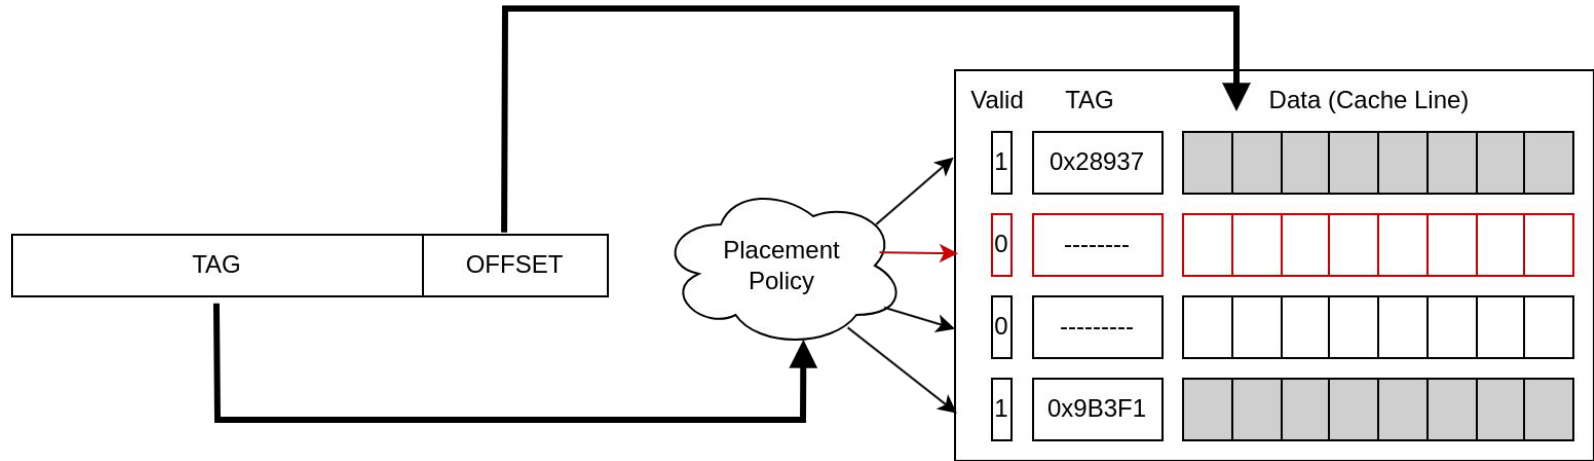
```
for(int i=0; i<256; ++i){
  for(int j=0; j<32; ++j){
    A[i][j] = A[i][j] + B[i][j]
  }
}
```

They will **never** hit the cache



Another approach to caches

- The static placement of direct-map under-utilizes the cache.
- Alternative: Place lines freely / dynamically on the cache → **Fully Associative Cache**
 - E.g. on the **next free line** (valid==0).



- More expensive, but ensures that the cache gets filled before evictions happen.

How do we find our data?

- Consider this **Fully Associative Cache** with all entries occupied:

New Access: 0x58732043

TAG OFFSET

Can this access be served
by the cache, or shall it go
to memory?

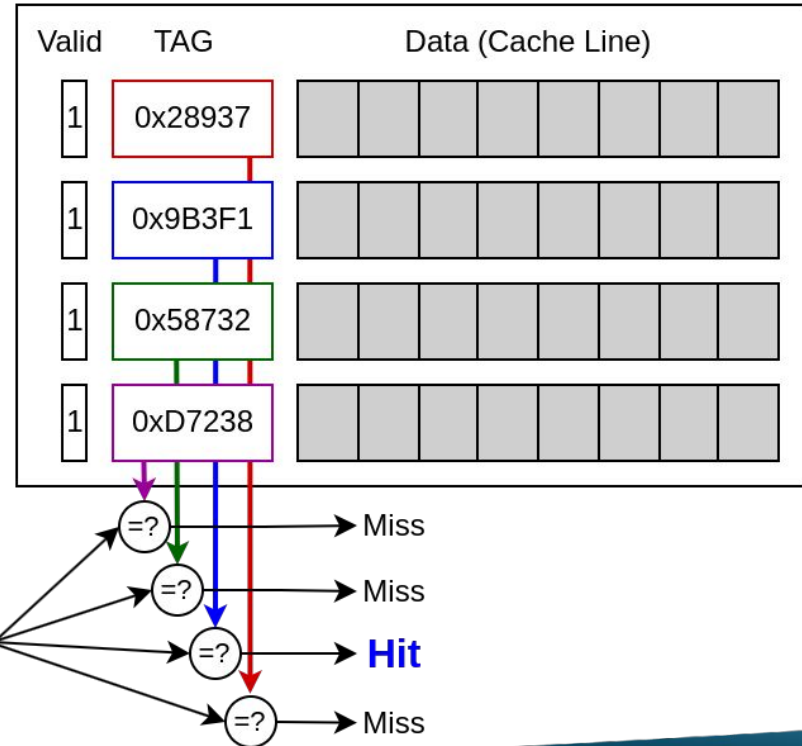
Valid	TAG	Data (Cache Line)
1	0x28937	
1	0x9B3F1	
1	0x58732	
1	0xD7238	

We look at the tags!

- We need to do as many comparisons as number of lines

New Access: 0x58732043

TAG OFFSET



Other considerations

- Assume this full cache again:

New Access: 0x4B94A022

TAG OFFSET

Does this access hit or miss?

Where do we put it? The cache is full...

Valid	TAG	Data (Cache Line)
1	0x28937	<div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>
1	0x9B3F1	<div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>
1	0x58732	<div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>
1	0xD7238	<div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>

Replacement Policies

- With direct mapping, we did not need replacement policies, as each address could only go to one specific line.
- Now, we have the freedom to choose who to evict. Some common policies are:
 - **FIFO**: First in - First out → Replace the oldest cache line
 - **LFU**: Least Frequently Used → Replace the cache line that got accessed the least
 - **LRU**: Least Recently Used → Replace the cache line that got used further in time
- There are many replacement policies, but some require adding extra logic, counters, and comparisons to the cache, will come at a cost

Complexity



“Good enough” policies

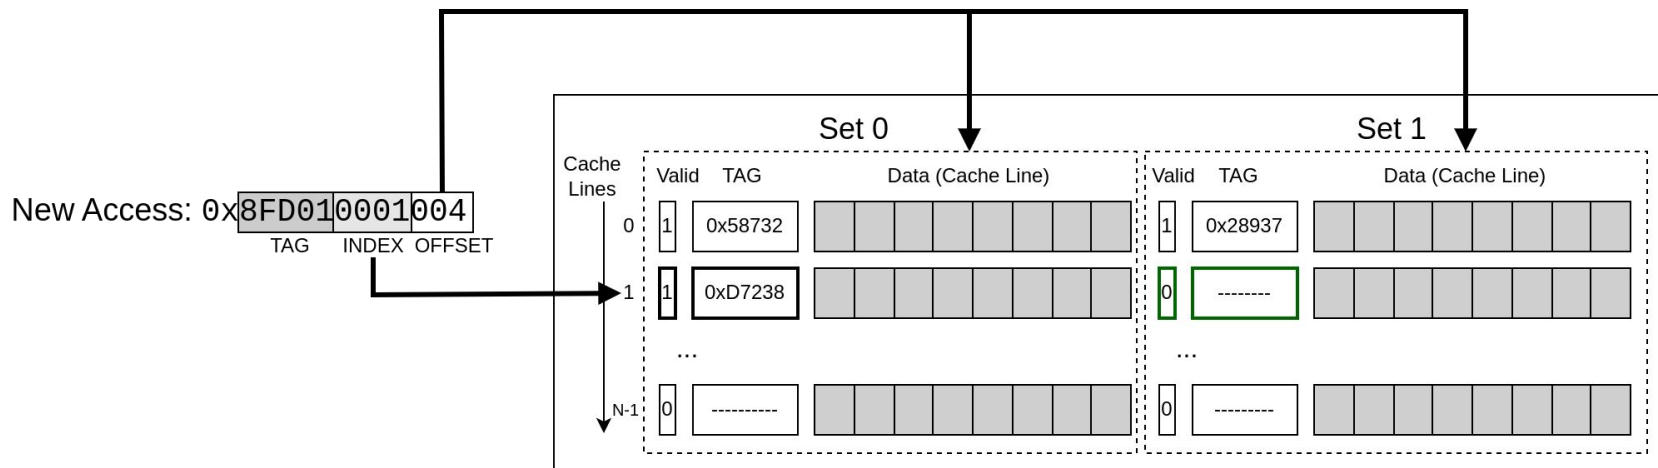
- Sometimes approximations to LRU / LFU are used
- For example: **Hot Potato**
 - 1. The first line of the cache is given a “Hot Potato”
 - 2. Whenever it is accessed, it passes the Hot Potato to the next line
 - 3. If the cache is full and we need to evict a line, we choose the one with the Hot Potato
- Very simple mechanism (you only need a counter)
- Lines that get used the most are less likely to hold the Hot Potato (but still possible)

Verdict on Fully Associative Caches

- Benefits
 - Ensures that the cache gets full before evictions happen.
 - Allows to control who gets evicted.
- Drawbacks
 - As many comparisons as Cache Lines (*1MB cache with 64B Lines → 16384 comparisons/access*)
 - Expensive in hardware → Limit on how big they can be.

The best of both worlds

- What if we mix direct mapping with associativity? → **Set Associative Cache**
- Addresses are mapped to cache lines, but can go into S sets
 - We refer to S-way Associative Cache (e.g. 2-way)



- Placement and replacement between sets is handled like in an associative cache.
- #S comparisons per access.

Verdict on Set Associative Caches

- Benefits

- Can avoid some collisions (e.g. This example from before)

```
int A[256][32] → 0xF3B44A57D6010000  
int B[256][32] → 0xF3B44A57D6020000  
for(int i=0; i<256; ++i){  
    for(int j=0; j<32; ++j){  
        A[i][j] = A[i][j] + B[i][j]  
    }  
}
```

} Even if they share a Line,
they can go to different sets

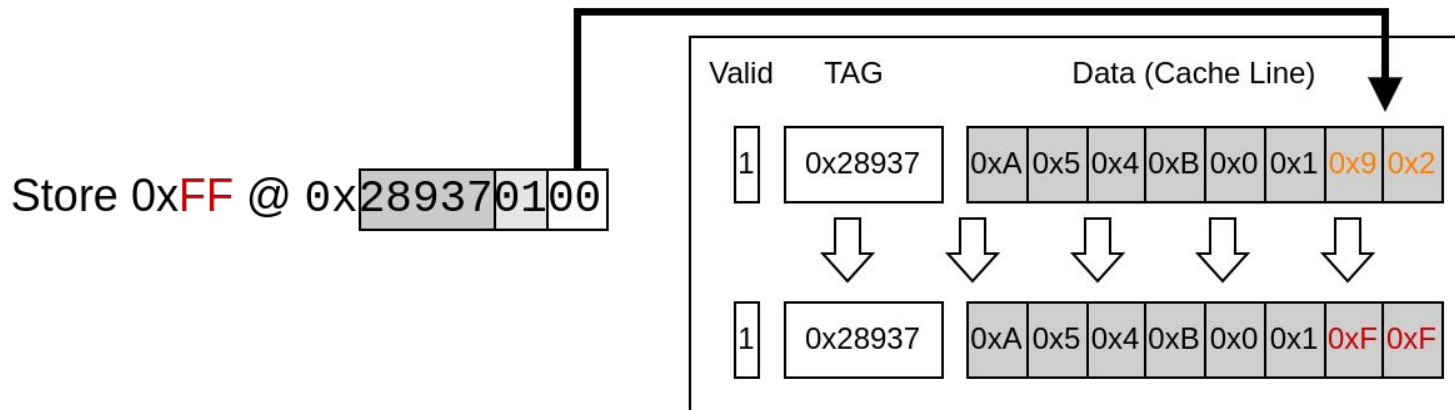
- Cheaper than Full Associative Caches (fewer comparisons and logic)

- Drawbacks

- “*Jack of all trades, master of none*” → More Costly than Direct Map, can’t avoid all collisions.
 - No optimal configuration (i.e. 2-way? 4-way? ... Depends on the workload!)

Modify data that is cached

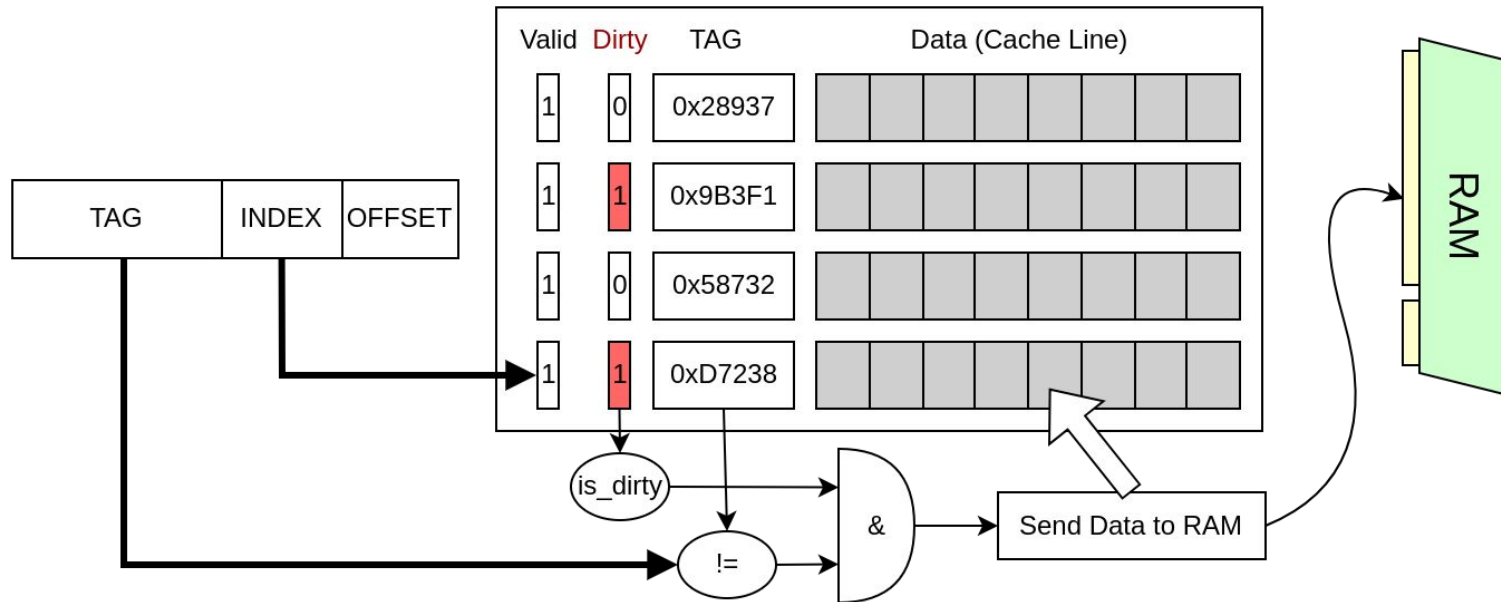
- What do we do if we want to **store** new data to a cached address?
- **First approach** → Modify the cache line:



What happens when this Line is evicted?

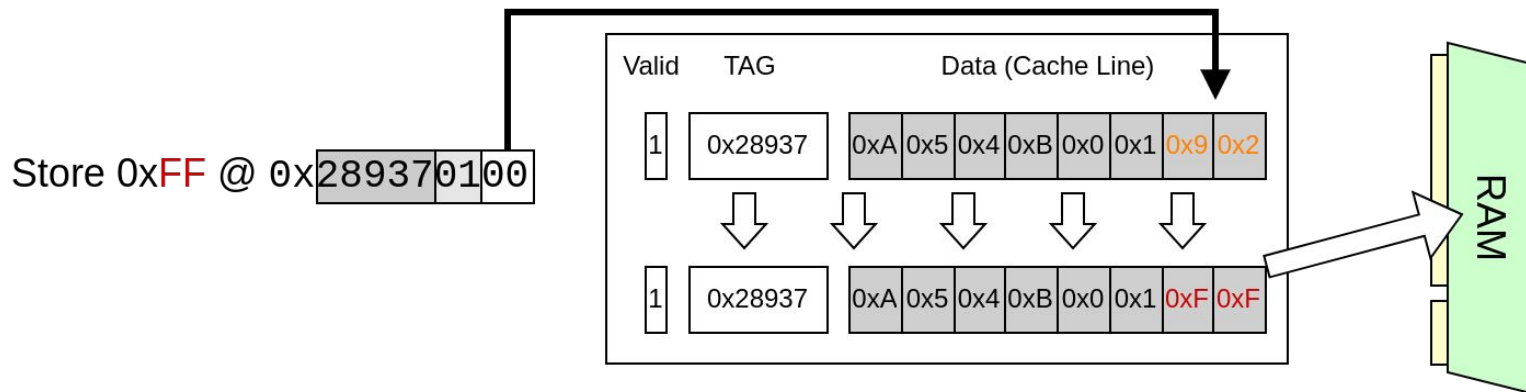
Propagate writes to the main memory

- We set a “**Dirty**” bit when we modify a Cache Line
- When we evict that line, we also **write-back** to the main memory



Another solution

- Some caches are “**write-through**” instead of “**write-back**”.
- When a Cache Line gets updated, so does the main memory




- More traffic to the main memory
- Higher latencies
- Defeats the purpose of the cache...
- ... But helps keeping coherence, and avoids the dirty bit

What to do with missing stores

- “**write-no-allocate**” policy:
 - Common policy in “**write-through**” architectures.
 - Missing writes directly modify the main memory without using the cache
 - Simple approach, with fewer evictions,
 - But increases main memory traffic, defeating the purpose of a cache
- “**write-allocate**” policy:
 - Common policy in “**write-back**” architectures.
 - Missing writes do not modify the main memory, only the cache.
 - Leaves the modification for later, when line is evicted
 - Can save up on main memory traffic if line is modified multiple times.
 - May evict a dirty line, which in turn means writing to the main memory anyway.

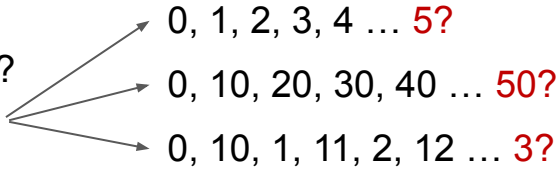
Can we improve cache hits further?

- We exploit the regular nature of most codes with spatial locality, but we still miss:



What's the next access? → Most likely **0x032678f024** !

Prefetchers

- We can leverage these “guesses”.
- The Load / Store unit can include a **Prefetcher**:
 - Analyzes the past accesses
 - Predicts future accesses, puts them in cache before they happen
- Varying complexity:
 - How far into the past can they look?
 - What patterns to they recognize?


0, 1, 2, 3, 4 ... 5?

0, 10, 20, 30, 40 ... 50?

0, 10, 1, 11, 2, 12 ... 3?
- Not always helpful:
 - Polluting the cache with speculative data can hurt the performance!


What about a second Cache?

- Modern architectures have different “Levels” of caches:
 - A first level (**L1**) inside the core, relatively **small** (16/32/64 KB) but **fast**
 - More levels (**L2, L3, ...**) outside the core, **bigger** (1MB - 100MB) but **slower**
- When an access misses on a cache level, it queries the next one.



Are caches only used to avoid accessing the RAM?

- No! But first, a 1-minute lecture on virtual memory.
- Loads/Stores can index addresses ranging from 0 to $2^{64}-1 \rightarrow$ The **address space**.
- The O.S. breaks the address space into **pages** of 4KB.
- Pages are, for example, used to load data from Disk into RAM.



Pages and virtual memory

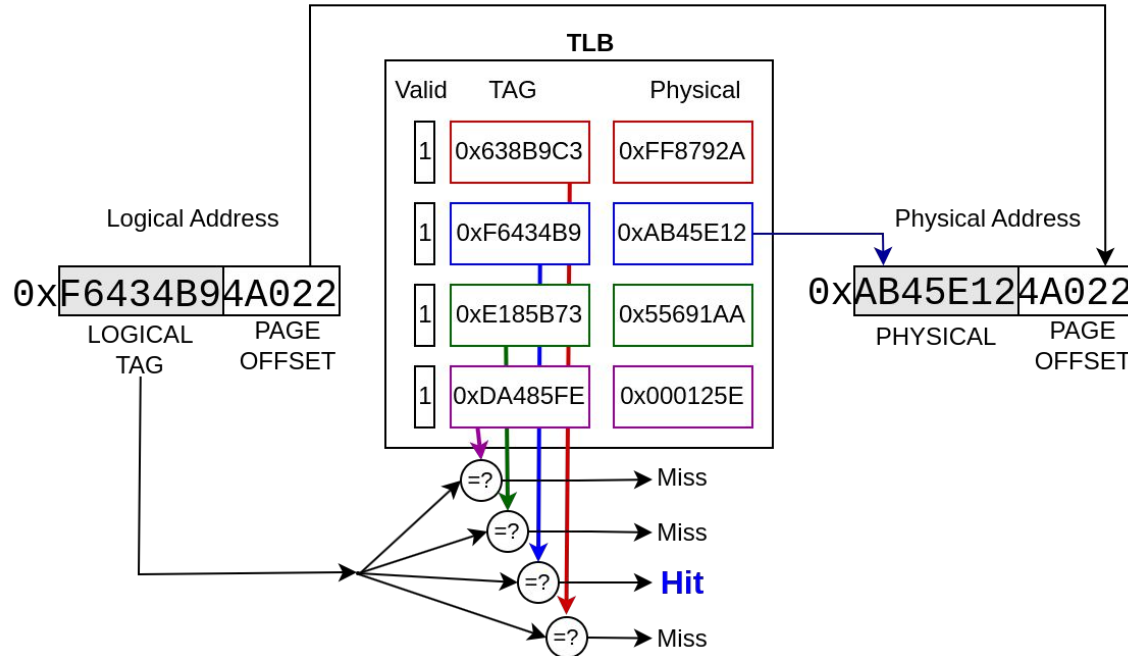
- Pages are also useful for **Virtual Memory**:
 - The O.S translates the addresses that users request to physical addresses (for security).
 - This is done in the **MMU** (Memory Management Unit) using a **Page Table**
 - Page table can be huge → And slow!

Page Table

Physical	Virtual
...	...
0x1A8	0x325
0x1A9	0xFF8
0x1AA	0x023
0x1AB	0x55E
0x1AC	0xFD1
0x1AD	0x123
...	...

Translation Lookaside Buffer (TLB)

- The hardware can help the O.S. by putting the page table in a cache → the **TLB**
- Whenever we need to translate an address from virtual to physical space, we use it:



- We want to avoid expensive misses, so we make this cache **Full-Associative** (and small!)
- On miss, access the page table.

Conclusions

- A fast memory access needs small and close storage (not RAM!)
- We can exploit access regularity:
 - Temporal locality → Reusing data
 - Spatial locality → Using data that is close together
- We create small memories (caches) to store frequently used data
 - Direct map, Associativity, write modes, replacement policies, ...
- These caches can be stacked into multiple levels
- They can be used on other contexts (e.g. address translation)